

Rajat Mohan Srivastava

AI-Enabled Data Analyst | GenAI & Agentic AI Systems

Noida, India • +91 88745 11637 • srajatmohan@gmail.com • [LinkedIn](#) • [Portfolio](#)

PROFESSIONAL SUMMARY

Data and applied AI analyst with 5 years at Ericsson, focused on building and shipping GenAI and agentic AI systems on enterprise telecom and operations data. Designs LangChain-based multi-agent workflows that combine LLM reasoning (OpenAI GPT, Google Gemini and Qwen models), retrieval-augmented generation (RAG) over vector databases, and orchestration with live KPI feeds and ML model outputs. Takes solutions from local prototype (Ollama on internal Eric AI infrastructure) to production on Google Cloud, packaging agents as Docker containers and deploying them on Vertex AI with Cloud Run, with data fetched directly from GCP pipelines (PMGCP). Brings a strong Python and SQL foundation from years of data engineering and automation, a cost-aware and privacy-first approach to model hosting, and a track record of partnering with product, operations and leadership teams to turn business questions into grounded, automated AI answers and executive-ready outputs.

KEY SKILLS & TECHNICAL SUMMARY

GenAI & Agentic AI: Agentic AI Systems, Multi-Agent Orchestration, LangChain (Agents, Chains, Memory), LLM Integration (OpenAI GPT, Google Gemini, Qwen; Ollama local deployment), RAG Pipeline Design, Vector Embeddings, Vector Databases, Semantic Search, Prompt Engineering, Chunking/Context Optimization, LLM Architectures & Fine-Tuning, Responsible & Ethical AI

Machine Learning: Supervised & Unsupervised Learning, Linear & Logistic Regression, Decision Trees, Random Forest, XGBoost, K-Means Clustering, Time-Series Forecasting (ARIMA, Prophet), Anomaly Detection, Feature Engineering, Model Evaluation (Precision/Recall, F1, RMSE), Scikit-learn, PyTorch (basics)

Cloud, Deployment & MLOps: Google Cloud Vertex AI, Cloud Run (serverless), Docker (containerization), Kubernetes, GCP data pipelines (PMGCP), Google BigQuery, CI/CD Pipelines, Model Lifecycle Management, Model Governance & Drift Monitoring, AWS & Azure AI Services, Cost-aware AI engineering

Programming & Data Engineering: Python (Pandas, NumPy, Scikit-learn, scripting & automation), SQL (MySQL, MS SQL Server, BigQuery), ETL/ELT design, Data pipelines, Data modeling, Data cleaning & validation, Technical specifications

Analytics & Visualization: Power BI (DAX, Power Query), SAP BI, KPI modeling, Automated PDF/PPT generation, Executive data storytelling

Methodologies & Tools: Agile/Scrum, Lean, ITIL4, Jira & Confluence, SharePoint data repositories, Data governance, QA & validation frameworks

PROFESSIONAL EXPERIENCE

AI-Enabled Data Analyst — Ericsson India Pvt Ltd

Noida, India | Feb 2021 – Present

- Lead the design and delivery of GenAI and agentic AI solutions for telecom operations, owning them end to end from requirement analysis and technical specification through development, testing, cloud deployment and support.
- Built a LangChain-based multi-agent KPI intelligence platform (OpenAI GPT, Gemini, Qwen, RAG, vector database) that answers complex business questions in natural language and auto-generates executive reports and presentations, significantly reducing manual analyst effort.
- Containerized the agentic solution with Docker and deployed it to production on Google Cloud Vertex AI using Cloud Run, integrating the PMGCP pipeline so agents fetch operational data directly at source.
- Engineered the RAG retrieval layer (vector embeddings, chunking, semantic search and context optimization) to ground LLM answers in historical operational data and reduce unsupported responses.
- Prototyped with locally hosted Qwen models (Ollama) on internal Eric AI infrastructure to keep data private and control inference cost, applying early governance and responsible-AI practices.
- Built Python (Pandas, NumPy, Scikit-learn) and SQL pipelines, with feature engineering and regression and time-series forecasting models on telecom and ticketing KPIs, feeding both the AI agents and regional and global dashboards.
- Automated reporting with Power BI and SAP BI for logistics, service KPIs and network performance, improving reporting turnaround by ~50%, with threshold and trend-based anomaly detection for network health.
- Led multi-layer data quality and validation checks, reducing reporting errors and defining metric SLAs with operations, product and leadership.
- Delivered 100+ analytics and automation projects across operations, logistics, finance and network teams.
- Partnered with product, operations and leadership to turn business needs into AI use cases and clear technical trade-offs.
- Documented BRDs, requirements and use cases in Jira/Confluence and delivered incrementally in Agile sprints; supported production pipelines using ITIL4 incident and change practices.

PROJECT HIGHLIGHTS

1. Agentic AI–Powered Telecom KPI Intelligence Platform (OpenAI GPT, Google Gemini, Qwen, Ollama, LangChain, RAG, Vector DB, Multi-Agent Orchestration, Google Vertex AI, Cloud Run, PMGCP Pipeline, Python, SQL, Power BI, PPT/PDF automation)

- Architected and shipped an agentic GenAI application that answers complex business KPI queries and automates executive deliverables (PDF reports, PPT presentations) for cross-functional teams, taking it from local prototype to cloud production deployment.
- Designed a Retrieval-Augmented Generation (RAG) pipeline using LangChain, vector embeddings and a vector database to retrieve and synthesize insights from large volumes of historical operational data for accurate, contextual responses.
- Orchestrated multi-agent workflows that combine the LLM with live KPI feeds and ML model outputs (Random Forest / XGBoost predictions and time-series forecasts), enabling joint historical and predictive analysis for root-cause analysis and trend forecasting.
- Applied prompt engineering, chunking and context optimization to improve answer relevance and grounding in source data.
- Prototyped on internal enterprise AI infrastructure (Eric AI) with locally hosted Qwen models (Ollama), keeping data private and avoiding per-token API cost during development, with early governance and responsible-AI controls built in.
- Deployed the agentic solution to production on Google Cloud Vertex AI, running it as a containerized, serverless service on Cloud Run for scalable, pay-per-use hosting.
- Integrated the agent directly with the PMGCP data pipeline on GCP to fetch operational KPI data at source, removing manual data extracts and keeping responses current.
- Automated executive-ready presentations, narrative summaries, dashboards and PDF exports, significantly reducing manual analyst effort.

2. HR & Finance Analytics Suite (Power BI, DAX, Power Query, SQL, Star-Schema Data Modeling, Row-Level Security)

- Designed and delivered an end-to-end Power BI analytics suite unifying HR and Finance data into a single source of truth for leadership, replacing fragmented spreadsheet-based reporting.
- Engineered a star-schema data model integrating HR (headcount, hiring, attrition, tenure) and Finance (budget, actuals, cost centers, operating expenses) datasets, with conformed date, department and cost-center dimensions for consistent cross-functional analysis.
- Built ETL and transformation logic in Power Query and SQL to cleanse, standardize and merge multi-source HR and finance extracts, with validation checks for reconciliation against source systems.
- Developed advanced DAX measures for workforce and financial KPIs, including attrition rate, time-to-hire, cost per employee, budget vs. actual variance, YoY/MoM trends and rolling forecasts.
- Implemented row-level security to protect sensitive employee and salary data, giving each stakeholder group access only to its own departments and cost centers.
- Created interactive executive dashboards with drill-through, bookmarks and what-if analysis, enabling leaders to move from organization-level trends to department-level root causes in a few clicks.
- Automated scheduled refresh and standardized KPI definitions with HR and Finance stakeholders, eliminating recurring manual month-end reporting and improving trust in the numbers.

3. ETL-Driven Inventory Analysis

- Built a Python-based ETL pipeline integrated with Excel Power Query to process inventory data from multiple sources, with automated data quality checks and incremental refresh.
- Impact: Reduced report delivery time from 1 day to ~30 minutes.

EDUCATION

B.Tech in Electrical & Electronics Engineering — Ajay Kumar Garg Engineering College, Ghaziabad | June 2019

CERTIFICATIONS

- AI Engineer Agentic Track: The Complete Agent & MCP Course – Udemy
- Building AI Agents using LangChain – Udemy
- ITIL4 Foundation Certified – Peoplecert International Limited
- Advanced Excel Formulas & Functions – Udemy
- SQL for Data Analysis: Advanced SQL Querying Techniques – Udemy